

COMMUNITY OVER CODE · GLASGOW 2026

WireGuard for Apache NuttX

A WireGuard VPN that works as a normal network device, `wg0` — on real boards, and inside a real NuttX kernel build.

ESP32-S3 & SPRESENSE

TUNNEL UP OVER REAL WI-FI

Real WireGuard peers

LINUX KERNEL, WINDOWS CLIENT

rv-virt knetnsh64

REAL KERNEL BUILD (BUILD_KERNEL)

NuttX 13.0.1 / master / 12.7.0

APPS V0.1.1 · ONE SOURCE

Satoru Akita · Sony Semiconductor Solutions



Agenda

Seven parts, about thirty minutes.

- | | | |
|---|---|---------|
| 1 | Why I needed a VPN on a NuttX device
What NuttX is, how this project started, and the problem I had to solve. | ~10 MIN |
| 2 | How I reused an existing WireGuard
Take wireguard-lwip and fit it into NuttX's own network device. | ~2 MIN |
| 3 | Bugs that a working demo hid
Five bugs that only showed up in deeper tests. | ~4 MIN |
| 4 | Moving the driver into the kernel
From an app to a real kernel device — and two more bugs. | ~2 MIN |
| 5 | How I checked that it really works
Seven test setups, what ping can't tell you, and one design choice. | ~10 MIN |
| 6 | Using a real board through the tunnel
Log in to a robot through the tunnel, and make it talk. | ~2 MIN |
| 7 | What comes next, and what to take home
Running it for real, other CPUs, sending it to NuttX, and three lessons. | ~3 MIN |

Agenda

Seven parts, about thirty minutes.


- | | | |
|---|---|---------|
| 1 | Why I needed a VPN on a NuttX device
What NuttX is, how this project started, and the problem I had to solve. | ~10 MIN |
| 2 | How I reused an existing WireGuard
Take wireguard-lwip and fit it into NuttX's own network device. | ~2 MIN |
| 3 | Bugs that a working demo hid
Five bugs that only showed up in deeper tests. | ~4 MIN |
| 4 | Moving the driver into the kernel
From an app to a real kernel device — and two more bugs. | ~2 MIN |
| 5 | How I checked that it really works
Seven test setups, what ping can't tell you, and one design choice. | ~10 MIN |
| 6 | Using a real board through the tunnel
Log in to a robot through the tunnel, and make it talk. | ~2 MIN |
| 7 | What comes next, and what to take home
Running it for real, other CPUs, sending it to NuttX, and three lessons. | ~3 MIN |

About Me

From Nagoya, through space, robots, and many conferences.



Satoru Akita — edge AI engineer at Sony Semiconductor Solutions, working with Apache NuttX from the application side.



UNIVERSITY

Space

Robotics

MEMS

Semiconductors



Many hackathons



Runs software conferences




Wrote a book on open source




Robotics

Molecular robotics, BIOMOD world champion (2015)




Computer Vision

AI and image processing systems



Edge AI

AITRIOS edge AI cameras



Embedded Systems

SPRESENSE, embedded app contest — world 3rd




NuttX / WireGuard

This talk

Overview: the project in one slide

This project adds WireGuard to NuttX as a network device.




Packets sent to `wg0` are encrypted and sent out as UDP. UDP that comes in is decrypted and handed over as if it came in on `wg0`. The other side can be any normal WireGuard peer.

TAKEAWAY Apps see `wg0` as a normal network interface. Only the "wire" under it is different: a UDP socket, not a cable.

How it started


From an idea in February to this talk in October.



TAKEAWAY The GSoC proposal was not selected, but the work went on. The kernel driver, the harder half, came last.

Background: the OS behind this work

What is Apache NuttX?



RTOS

POSIX APIs

BSD sockets

Filesystem / VFS

NSH shell

Built-in networking

Device drivers

Apache top-level project

TAKEAWAY NuttX lets you program small embedded devices the Unix way: POSIX APIs, BSD sockets, files, and devices.

Background: NuttX and other RTOSes

Where NuttX fits — and where it is used.

RTOS	MAIN FOCUS	TYPICAL USE	ARCHITECTURES	BOARDS
FreeRTOS	Small real-time kernel	MCU firmware, control, IoT	15+ CPU architectures	300+ supported boards
Zephyr	All-in-one platform for connected devices	Wearables, industrial & connected IoT		
NuttX	Unix-like embedded OS	Networking, file systems, drones, sensing	CONFIGS 1,500+ configuration templates	IN PRODUCTION Sony · PX4 drones · Xiaomi OpenVela (1,000+ SKUs) · Japan's 2024 Moon mission

FreeRTOS/FreeRTOS 7.8k★ · zephyrproject-rtos/zephyr 16.5k★ · apache/nuttX 4.1k★

Stars show how many developers notice a project, not how many devices use it.

TAKEAWAY NuttX has fewer stars than FreeRTOS and Zephyr, but it already runs in real products, from drones to the Moon.

Sources: nuttx.apache.org · nuttx.apache.org/market-share · docs.px4.io · GitHub star counts, October 2026

History: a long road with NuttX

NuttX at Sony — from audio to edge sensing.


- 2015**
 - **Sony audio products**
NuttX inside Sony audio products.
- 2018–19**
 - **SPRESENSE / CXD56xx**
Audio, GNSS, camera and multi-core support added to NuttX; shown at Open Source Summit + ELC Europe 2019.
- 2020**
 - **Upstream collaboration**
Sony engineers become major NuttX contributors: “A Journey from Fork to Mainline,” NuttX Online Workshop 2020.
- 2020s**
 - **Sensing & edge devices**
More camera, GNSS and sensing products.

TAKEAWAY Sony used NuttX long before edge AI. Years of audio products and SPRESENSE led to sensing and camera systems.

Sources: M. Ishikawa (Sony), “NuttX for Embedded Linux Developers,” OSS+ELC Europe 2019 · A. Jerpelea (Sony), “A Journey from Fork to Mainline,” NuttX Online Workshop 2020 · nuttx.apache.org/docs (CXD56xx)


Experience: where I come from

NuttX in SPRESENSE and AITRIOS — my part.



SPRESENSE

NuttX SDK · sensing & camera · GNSS / LTE · prototyping and embedded development.



AITRIOS / edge AI camera


Edge AI · setting up cameras · ESP32 + NuttX · running them on site, not on a desk.

Where I work: application development → system integration → deployment & operation.

TAKEAWAY My problem was not making NuttX run. It was reaching NuttX devices after they were installed.

Problem: Wi-Fi security stops too early

Reaching a device from far away is hard.



TAKEAWAY What was missing: a standard, secure path all the way to the NuttX device.

Options: MQTT, HTTP, router VPN, ...

Other ways to reach a device — compared.


WAY	WHO CONNECTS	ANY SERVICE (TELNET, WEB...)?	SERVER IN THE MIDDLE?	SECURITY ENDS AT	ON NUTTX TODAY
Open a port to the Internet	engineer → device	yes	no	the app (often nothing)	easy, but open to everyone
HTTPS to a cloud API	device → cloud (polling)	no, only your API	yes, a web server	the cloud server	<code>webclient</code> , <code>libcurl4nx</code> + <code>mbedtls</code>
MQTT broker (TLS)	device → broker	no, messages only	yes, a broker	the broker	<code>mqttc</code> , <code>paho_mqtt</code> + <code>mbedtls</code>
Reverse tunnel / relay (SSH, WebSocket)	device → relay	one port per tunnel	yes, a relay host	the relay host	<code>dropbear</code> , <code>libwebsockets</code>
VPN on the router	router ↔ office	yes, the whole LAN	no	the router; the LAN hop is open	no change on the device, but you need that router
Carrier closed network (private APN, SIM VPN)	through the carrier	yes	yes, the carrier	the carrier gateway	LTE modem; cellular only
VPN on the device (this talk)	either side	yes, any IP service	no, any WireGuard peer	the device itself	<code>wg0</code>

TAKEAWAY MQTT and HTTPS are great for data. To log in and fix a device, end to end, with no server in the middle, the VPN has to be on the device.

Mesh VPN agents (Tailscale, ZeroTier) are another “VPN on the device” option, but they usually need a bigger OS. NuttX clients: `nuttX-apps 13.0.1` `netutils/` (`mqttc`, `paho_mqtt`, `webclient`, `libcurl4nx`, `libwebsockets`, `dropbear`) and `crypto/mbedtls`.

Solution: a VPN on the device

Why WireGuard is a good fit.



- **Modern IP-layer VPN**

small and simple: much less code than IPsec or OpenVPN.

- **UDP transport, public-key peers**

no certificates and no session setup to manage.


- **Already everywhere**

in the Linux kernel, with an official Windows client. It works with existing peers, not a custom protocol.

TAKEAWAY WireGuard solves the access problem with a standard protocol, not a homemade security layer or a vendor cloud.

Architecture: where wg0 sits

WireGuard on the NuttX network stack.



TAKEAWAY To NuttX, wg0 is a normal network device. Under it, WireGuard turns its packets into encrypted UDP.

Agenda

Seven parts, about thirty minutes.

- 1 Why I needed a VPN on a NuttX device** ~10 MIN
What NuttX is, how this project started, and the problem I had to solve.
- 2 How I reused an existing WireGuard** ~2 MIN
Take wireguard-lwip and fit it into NuttX's own network device.
- 3 Bugs that a working demo hid** ~4 MIN
Five bugs that only showed up in deeper tests.
- 4 Moving the driver into the kernel** ~2 MIN
From an app to a real kernel device — and two more bugs.
- 5 How I checked that it really works** ~10 MIN
Seven test setups, what ping can't tell you, and one design choice.
- 6 Using a real board through the tunnel** ~2 MIN
Log in to a robot through the tunnel, and make it talk.
- 7 What comes next, and what to take home** ~3 MIN
Running it for real, other CPUs, sending it to NuttX, and three lessons.

Porting: four functions to write

Reusing wireguard-lwip without lwIP.

PROTOCOL & CRYPTO — KEEP

Handshake, transport, BLAKE2s, ChaCha20-Poly1305, Curve25519. Portable C, no OS dependency. Used as-is.

PLATFORM HOOKS — ADAPT

Everything OS-specific sits behind `wireguard-platform.h` — four functions: clock, random bytes, TAI64N time, load check.

LWIP NETWORK GLUE — REPLACE

`wireguardif.c` 's lwIP `netif` registration and buffers. NuttX has no lwIP, so this layer is the real porting work.

TAKEAWAY A clean line between the protocol and the OS let the protocol and crypto move unchanged. Only the OS and network glue had to change.

Design: one virtual network card

Mapping lwIP onto NuttX netdev.

CONCEPT	LWIP	NUTTX
Network device	struct netif	struct net_driver_s
Buffers	pbuf	iob
Registration	netif_add()	netdev_register()
Start sending	callback	devif_poll()

`wg0` is a `NET_LL_TUN` netdev, and its wire is a UDP socket.

TAKEAWAY Porting meant using NuttX's own building blocks, not bringing lwIP into NuttX.

Result: the first working tunnel

A working tunnel is only the beginning.

LAYER	WORK NEEDED	IN THIS PORT
Protocol & crypto	None — portable C, as-is	3,079 lines, exactly the same as upstream
OS hooks	Small — four functions	clock, /dev/urandom, TAI64N, load
Network glue	The real work	netdev + UDP socket + RX task

It worked: `wg genkey / set / setconf` **at run time. No rebuild to change keys or peers.**

And this is where the interesting part starts.

TAKEAWAY Getting a tunnel up was the easy step. Almost everything interesting came after it already “worked.”


Agenda

Seven parts, about thirty minutes.

- | | | |
|---|---|---------|
| 1 | Why I needed a VPN on a NuttX device
What NuttX is, how this project started, and the problem I had to solve. | ~10 MIN |
| 2 | How I reused an existing WireGuard
Take wireguard-lwip and fit it into NuttX's own network device. | ~2 MIN |
| 3 | Bugs that a working demo hid
Five bugs that only showed up in deeper tests. | ~4 MIN |
| 4 | Moving the driver into the kernel
From an app to a real kernel device — and two more bugs. | ~2 MIN |
| 5 | How I checked that it really works
Seven test setups, what ping can't tell you, and one design choice. | ~10 MIN |
| 6 | Using a real board through the tunnel
Log in to a robot through the tunnel, and make it talk. | ~2 MIN |
| 7 | What comes next, and what to take home
Running it for real, other CPUs, sending it to NuttX, and three lessons. | ~3 MIN |

The Testing Pattern

Simple tests pass. Deeper tests fail.



A simple test passes. A deeper one fails.

Because the simple test passed, you trust that part and spend hours looking somewhere else. This happened again and again in this project, and not only in the apps version.

Bug 1 / 7

A return code of 0 that did nothing.

`setsockopt(SO_RCVTIMEO)` returns **0 (success)**.

But `recvfrom()` still **blocks forever**. The timeout was accepted, then ignored.

A perfect deadlock:

timer never fires → no handshake is sent → the peer stays silent → `recvfrom()` never returns.

TAKEAWAY `setsockopt()` returning 0 does not prove the option works. An unsupported option can show up as a deadlock, not as a clear error.

Bug 2 / 7

Created successfully. Did not survive.

The RX thread was started with `pthread_create()` + `pthread_detach()`.

When the `wg` command exited, the thread **stopped** (not even visible in `ps`).

A `usleep()` right after create made the thread's `printf` appear, so it **was** created.

My mistake: I thought `detach` meant “lives longer than the command.” It only controls **when the thread's memory is freed**, not whether it keeps running after the command exits.

TAKEAWAY `pthread_detach()` decides when memory is freed, not how long the thread lives. A detached thread still dies with the task group that started it.

Bug 3 / 7 · my favorite

ping works. TCP dies.

Same tunnel: `ping` is **0% loss**. A telnet 3-way handshake even **completes**, then the connection dies.

```
wg_encrypt_and_send: psock_sendto  
ret=-1 errno=9 EBADF
```

The sends that **worked** ran inside the `wg_rx` task. Sends from another task got EBADF.

Cause: an fd belongs to one **task group**; it is not global.

Fix: use the internal API without fds: `psock_*()` on a `struct socket`, which any task can use.

TAKEAWAY In NuttX, a file descriptor belongs to one task group. Another task cannot use it, even though it points to the same socket.

Bug 4 / 7

“No real impact” was 10×.

The fix without fds also broke `poll()` . Quick fix: check every 50 ms.

I thought: “it only wakes the CPU when idle.” Wrong: each packet waited up to 50 ms, and that made it **10× slower**.

Fix: hook `psock_poll()` with a callback and a semaphore (a real blocking wait).

TAKEAWAY The 50 ms polling kept things correct but made it about 10× slower. A quick fix you never measure can look free and cost 10×.

Bug 5 / 7

sim passes. Real hardware is different.

Saving the config with

`wg showconf > /data/wg0.conf` froze the board,
only on **real hardware with SPIFFS**, never in sim.
Fixed with `wg saveconf`, which writes the file itself.

The `wg_rx` task calls `ipv4_input()` on **its own stack**, so the whole TCP call chain runs on it, plus encrypting the reply.

Default **3072 B** was a number nobody had measured. On the board it was **94% used, 176 B left** (`CONFIG_STACK_COLORATION` + `ps`).

TAKEAWAY A default stack size is only a guess until you measure it on real hardware. Sim's big stacks hide the problem.

Agenda

Seven parts, about thirty minutes.

- | | | |
|---|---|---------|
| 1 | Why I needed a VPN on a NuttX device
What NuttX is, how this project started, and the problem I had to solve. | ~10 MIN |
| 2 | How I reused an existing WireGuard
Take wireguard-lwip and fit it into NuttX's own network device. | ~2 MIN |
| 3 | Bugs that a working demo hid
Five bugs that only showed up in deeper tests. | ~4 MIN |
| 4 | Moving the driver into the kernel
From an app to a real kernel device — and two more bugs. | ~2 MIN |
| 5 | How I checked that it really works
Seven test setups, what ping can't tell you, and one design choice. | ~10 MIN |
| 6 | Using a real board through the tunnel
Log in to a robot through the tunnel, and make it talk. | ~2 MIN |
| 7 | What comes next, and what to take home
Running it for real, other CPUs, sending it to NuttX, and three lessons. | ~3 MIN |

Part 4 • Going deeper

Into the kernel.

The limit of version 1

It lives in `apps/` and uses `psock_*` internals. Fine for a FLAT build. But it can never be a proper in-tree driver, and it does not work in **PROTECTED / KERNEL** builds, where apps and kernel are kept apart.

The move

Put the device **in the kernel** (`drivers/net/wireguard/`) and control it from a small user-space `wg` command with **ioctl**.

TAKEAWAY An app-level port proves the protocol works. Only a kernel driver is a proper device that also works in PROTECTED and KERNEL builds.

The Kernel Design

A kernel netdev, an ioctl ABI, a write-only key.

- **Kernel netdev**

`netdev_lowerhalf` with its own UDP socket, receive thread, timers and cookies. The driver owns the socket and the thread, so they no longer die with the command that started them.

- **Simple ioctl ABI with no pointers**

`SIOCS/GWGIF`, `SIOCS/D/GWGPEER`. No pointers, because in PROTECTED / KERNEL builds the kernel copies the caller's struct as it is.

- **Private key is write-only**

the get ioctl never returns the private key.

- **Kernel crypto is NuttX's own**

BLAKE2s, ChaCha20-Poly1305, Curve25519 from `crypto/`. No copied code in the kernel. The only copied file, an MIT-licensed X25519, is in the user-space command for offline `genkey`. Using the kernel's Curve25519 through cryptodev would pull in the whole software cipher set for one multiply.

TAKEAWAY The socket, the thread and the no-pointer ABI follow one rule: don't depend on code that can exit, or on memory the kernel can't reach.

Two more bugs: same shape

Deeper level, same kind of bug.

A bug in NuttX's own crypto

`chacha20poly1305` put the counter in the wrong nonce bytes. Counter 0 matches, so **the handshake succeeds and every data packet after the first fails**. Nothing in NuttX had ever used it. Two-line fix, going upstream as a separate `crypto:` PR.

A BUILD_KERNEL-only crash

`wg_set_if()` copied the peers into a **6 KB array on the stack**; the kernel stack is 3 KB → overflow → heap corruption → panic. **sim (FLAT, big stacks) never hit it**. Moved to the heap.

TAKEAWAY A working handshake does not prove the data path or the kernel code is right. Both bugs only showed up after the handshake, in a real kernel build.

Agenda

Seven parts, about thirty minutes.

- | | | |
|---|---|---------|
| 1 | Why I needed a VPN on a NuttX device
What NuttX is, how this project started, and the problem I had to solve. | ~10 MIN |
| 2 | How I reused an existing WireGuard
Take wireguard-lwip and fit it into NuttX's own network device. | ~2 MIN |
| 3 | Bugs that a working demo hid
Five bugs that only showed up in deeper tests. | ~4 MIN |
| 4 | Moving the driver into the kernel
From an app to a real kernel device — and two more bugs. | ~2 MIN |
| 5 | How I checked that it really works
Seven test setups, what ping can't tell you, and one design choice. | ~10 MIN |
| 6 | Using a real board through the tunnel
Log in to a robot through the tunnel, and make it talk. | ~2 MIN |
| 7 | What comes next, and what to take home
Running it for real, other CPUs, sending it to NuttX, and three lessons. | ~3 MIN |

Testing: simulator, QEMU, real boards

Seven places it must work.

WHERE	VERSION	WHAT I CHECKED
sim:wireguard	kernel driver	builds with no errors (the config CI builds)
rv-virt knetnsh64	kernel · BUILD_KERNEL	virtio-net tunnel to Linux; <code>wg</code> is a separate program that calls in through syscalls
rv-virt pnsh64	kernel · BUILD_PROTECTED	tunnel with the MPU split on: <code>wg</code> in <code>nuttx_user</code>, every ioctl crosses the boundary; tunnel works, but 2 of 8 repeated runs hit a fault (open)
rv-virt knetnsh64_smp	kernel · SMP, 4 CPUs	heavy traffic while changing the config and doing <code>down</code> / <code>up</code>; no fault, tunnel came back
SPRESENSE, measured	kernel · Cortex-M4F	under load and 40 <code>down</code> / <code>up</code> cycles: heap, buffers and stacks exactly the same; receive thread uses 1472 B of 6096
ESP32-S3 + SPRESENSE	kernel · real Wi-Fi	tunnel to the Windows client: ESP32 native 4/4, SPRESENSE via usrsock 6/6
same two boards	apps v0.1.1	Linux & Windows peers; telnet, HTTP, 7 MB, rekey, power-cycle

Both boards, PROTECTED and KERNEL, many CPUs: measured, not guessed.

TAKEAWAY A tunnel in sim is not enough. Test real builds, many CPUs, and real hardware.

Hardware tests used the version before the final rebase onto upstream; KERNEL/SMP tunnels were run again on the rebased version.

Testing: what ping misses

A working tunnel does not prove enough.

PROBLEM	WOULD PING NOTICE?	WHAT I MEASURED
The same private key on every boot: the <code>xorshift128</code> <code>/dev/urandom</code> backend can be seeded from build-time values	no, every normal test still passes	3 restarts → 3 different keys; a basic check of the entropy pool path, not proof of good randomness
Session keys left in memory after down	no	read the running memory: transport keys 32/32 bytes set while up → all zero after <code>down</code> ; exactly one copy of the static key, before and after
A leak that only shows after hours	no	151 forced rekeys past the key lifetime, 10 endpoint changes, 100 <code>down</code> / <code>up</code> : no buffers lost, allocations back to the boot count exactly, RX-thread stack peak 2.5 KB and not growing

Each of these passes every normal test while being broken.

TAKEAWAY A working tunnel does not prove unique keys, cleared secrets, or stable memory. Each needs its own check.

Design choice: trust the clock?

The timestamp problem has no easy answer.

A peer drops any handshake whose timestamp is not newer than the newest it has seen. So the device must always send a newer one, **even after a reboot**.

OPTION	FLASH WRITES	WORKS WITHOUT RTC	RESULT
Real-time clock + highest value this boot	none	no	chosen
Save every timestamp	every handshake	yes	wears out flash, slows the handshake
Reserve a range in flash	one per N	yes	needs a storage guarantee NuttX does not have
Wait until it catches up	none	no	you wait as long as it last ran

Same board, one change, nothing sent to it: **clock unset** → **no handshake in 75 s**; **clock set** → **4.1 s**.

Used the same rule as Linux and wireguard-go, and made the case it misses **print a warning.**

TAKEAWAY No clock, no newer timestamp: a device may be unable to connect out and still pass a ping test


Agenda

Seven parts, about thirty minutes.

- | | | |
|---|---|---------|
| 1 | Why I needed a VPN on a NuttX device
What NuttX is, how this project started, and the problem I had to solve. | ~10 MIN |
| 2 | How I reused an existing WireGuard
Take wireguard-lwip and fit it into NuttX's own network device. | ~2 MIN |
| 3 | Bugs that a working demo hid
Five bugs that only showed up in deeper tests. | ~4 MIN |
| 4 | Moving the driver into the kernel
From an app to a real kernel device — and two more bugs. | ~2 MIN |
| 5 | How I checked that it really works
Seven test setups, what ping can't tell you, and one design choice. | ~10 MIN |
| 6 | Using a real board through the tunnel
Log in to a robot through the tunnel, and make it talk. | ~2 MIN |
| 7 | What comes next, and what to take home
Running it for real, other CPUs, sending it to NuttX, and three lessons. | ~3 MIN |

Demo setup: the network

One tunnel, on a phone's local network.



TAKEAWAY The phone and the venue only ever see encrypted UDP, with no server in the middle and no cloud.

Demo: live on the StackChan

Log in through the tunnel, and talk to the robot.

```
1  wg show                # handshake, bytes
2  telnet 10.10.0.2        # log in
3  ifconfig                # wlan0 and wg0
4  stackchan face happy
5  stackchan say hello.wav
6  tcpdump: outside vs. inside
7  wg show                 # bytes went up
```

Outside the tunnel (the phone's network): only encrypted UDP on port 51820.

Inside the tunnel (wg0): the telnet text you just typed.

Backup video (earlier demo):
youtu.be/1kyX2av5WG4

Agenda

Seven parts, about thirty minutes.

- | | | |
|---|---|---------|
| 1 | Why I needed a VPN on a NuttX device
What NuttX is, how this project started, and the problem I had to solve. | ~10 MIN |
| 2 | How I reused an existing WireGuard
Take wireguard-lwip and fit it into NuttX's own network device. | ~2 MIN |
| 3 | Bugs that a working demo hid
Five bugs that only showed up in deeper tests. | ~4 MIN |
| 4 | Moving the driver into the kernel
From an app to a real kernel device — and two more bugs. | ~2 MIN |
| 5 | How I checked that it really works
Seven test setups, what ping can't tell you, and one design choice. | ~10 MIN |
| 6 | Using a real board through the tunnel
Log in to a robot through the tunnel, and make it talk. | ~2 MIN |
| 7 | What comes next, and what to take home
Running it for real, other CPUs, sending it to NuttX, and three lessons. | ~3 MIN |

Running it for real

From a demo to something you can run.

	BEFORE	AFTER
Keys	built into Kconfig; rebuild to change	genkey / set / setconf at run time; never in the firmware image
When it stops	guessing	check ROUTER / LAN / TUN / TCP every minute to find which part failed

Four separate checks answer one question: when the tunnel goes quiet, is it my side?

TAKEAWAY Run-time keys and per-layer checks turn "the tunnel is down" into "the router side is down." That is the difference between a demo and something you can leave running.

Other CPUs: beyond the first board

Did the “four functions” bet work?

The apps (FLAT) version ran on four CPU types with **no code change**:

TARGET	CODE CHANGE	RESULT
x86_64 (sim)	none	tunnel
ARM Cortex-A7 (QEMU)	none	tunnel
Xtensa LX7 (ESP32-S3)	none	tunnel over real Wi-Fi
ARM Cortex-M4F (SPRESENSE)	none	tunnel over real Wi-Fi (GS2200M)

TAKEAWAY Four small OS functions let one source run on four CPU types, unchanged. The bet paid off.

Checked with the apps (FLAT) version; the kernel driver was also checked on sim, rv-virt (BUILD_KERNEL) and both boards over real Wi-Fi (ESP32-S3 native, SPRESENSE over usisock).

Giving back: sending it to NuttX

The Apache Way: small steps, in the open.

- **A small crypto: PR first**

the ChaCha20-Poly1305 nonce fix, which helps NuttX on its own.

- **Then the driver PR**

`drivers/net/wireguard` — the kernel device, ABI, and its documentation.

- **Then the command PR**

`apps/system/wg` — the command that talks to the driver.

- **Design ready for `dev@nuttx.apache.org`**

passes checkpatch / nxstyle, split into PRs, licenses checked. Community over code.

TAKEAWAY Sending the crypto fix as its own PR lets NuttX take it even if the driver is not accepted.

Status: the work is split into small PRs that pass checkpatch and build one commit at a time; the dev@ discussion and the PRs come next.

Takeaways: what to take home

Three things to remember.

- **Simple tests pass, deep tests fail**

test deeply (long data runs, real hardware, a real kernel build), not just “the handshake worked.”

- **A clean OS boundary pays off**

four CPU types, three NuttX versions, no changes.

- **Porting to a project can make the project better**

two hidden NuttX bugs found on the way: its ChaCha20-Poly1305 nonce, and a boot bug that stops a clean release from reaching the shell on this board.

COMMUNITY OVER CODE · GLASGOW 2026

Thank you

Questions welcome — the kernel driver, the bugs, the ioctl ABI, or the hardware.

youtu.be/1kyX2av5WG4

DEMO

[apache/nuttx](#) · [apache/nuttx-apps](#)

GOING UPSTREAM

Satoru Akita · Sony Semiconductor Solutions

