

Speaking script — WireGuard for Apache NuttX

(Community Over Code, Glasgow)

Deck: [coc-glasgow-slides.html](#) (and the version on the official template, [coc-glasgow-reveal/](#)), 42 slides, English. This file is **made by** [export_en_script.py](#) from the speaker notes in the deck (press **N** there to see them). The "Say" text below *is* those notes. Do not edit this file by hand. Change the notes in the deck and run the script again, or the two will stop matching. The Japanese [practice script](#) uses the same slide order.

Total: about 36:40, including a 90-second demo. The times add up from the start ("you should be at about this point"), at 150 words per minute. This is longer than a 30-minute slot, so practice to finish the main talk in 27–28 minutes, and keep the cut list below ready.

The most important part is slides 31–33 (testing). Together they take 10.7 minutes of the talk, and they hold the claims that a reviewer will really check. Do not make them shorter to save time somewhere else. Cut before them or after them, never inside. They are one story: the first is where it runs, the second is what "it runs" does not prove, and the third is the one place where I had to decide instead of test. If you split them, the story breaks.

Do not cut slides 4–14 either (About Me to Architecture). What NuttX is, where it is used, Sony's history with it, how I met it, and why that led to WireGuard: everything after that depends on this.

If time is short, cut in this order:

1. Slide 40 (Giving Back). It helps the talk, but the talk does not depend on it. If you cut it, answer it in Q&A if someone asks.
2. Slide 39 (Other CPUs, the portability table). It adds support, but it is not the main story.
3. Slide 38 (Running It for Real).

If the live demo stops, switch to the recording right away. **The agenda slides (3, 15, 19, 26, 30, 34, 37) are a ten-second pause each:** point at the highlighted line and pause. Do not read the list. If you have no time at all, skip them without a word. They save almost no time, so cut the three slides above instead.

Check before the talk: the GitHub star numbers on slide 8 change. Check them again on the day, or keep the date on the slide.

1 — Title (0:00 → 0:43)

On slide: "WireGuard for Apache NuttX." ESP32-S3 & SPRESENSE tunnel over real Wi-Fi · real WireGuard peers (Linux kernel, Windows client) · runs in a real kernel build (BUILD_KERNEL) · one source across NuttX 13.0.1 / master / 12.7.0.

Say:

- Today I will talk about adding WireGuard to Apache NuttX.

- It is a WireGuard VPN that works as a network device called wg0. The tunnel runs over real Wi-Fi on ESP32-S3 and on SPRESENSE, and the other side is always a real WireGuard program — the Linux kernel module and the official Windows client. The apps version builds from one source on three NuttX versions. Later in the talk, you will see the new driver running inside real KERNEL and PROTECTED builds. That is the harder and more interesting half.
- This talk is about what “it works” was hiding, and what moving it into the kernel brought out.

2 — Agenda (0:43 → 1:17)

On slide: The seven parts as sentences with their times: why I needed a VPN on a NuttX device (~10 min) / how I reused an existing WireGuard (~2) / bugs that a working demo hid (~4) / moving the driver into the kernel (~2) / how I checked that it really works (~10) / using a real board through the tunnel (~2) / what comes next (~3).

Say:

- Here is the plan for this talk.
- There are seven parts. First, why I needed a VPN on a NuttX device: what NuttX is, how this project started, and the problem I had to solve. Second, how I reused an existing WireGuard. Third, five bugs that a working demo hid. Fourth, moving the driver into the kernel. Fifth, how I checked that it really works. Sixth, a short demo. And last, what comes next, and what to take home. About thirty minutes, then questions.

3 — Agenda: ① Why I Needed a VPN (1:17 → 1:27)

Say (one line): Let's start with part one.

Pause: point at the highlighted line, one beat.

4 — About Me (1:27 → 2:06)

On slide: Photo, and a timeline Robotics → Computer Vision → Edge AI → Embedded Systems → NuttX / WireGuard with a picture per step.

Say:

- First, a little about me.
- I am an edge AI engineer at Sony Semiconductor Solutions. I work with Apache NuttX from the application side. I was born in Nagoya. I went to university in Sendai, where I studied space, robots, MEMS and semiconductors. Now I live in Kawasaki. I often join hackathons, I help run software conferences, and I wrote a book about open source.
- My road to this project went through molecular robots, computer vision and AI systems, edge AI cameras, and embedded systems on SPRESENSE. NuttX and WireGuard are the latest stop on that road.

5 — Project Overview (2:06 → 2:46)

On slide: The wg0 netdev picture: application / NSH, the NuttX network stack, wg0 next to eth0 / wlan0, a UDP socket underneath, the Internet, the peer.

Say:

- Here is the whole idea in one picture.
- Apache NuttX is an RTOS that follows POSIX and has its own TCP/IP stack. It had no VPN. This project adds wg0 as a virtual network interface. Applications see a normal interface. Under it, the encrypted data travels over a UDP socket. The other side is a real WireGuard peer — Linux, Windows, or another gateway.
- Because NuttX follows POSIX and uses BSD sockets, it is a good fit for this work. FreeRTOS and Zephyr can run a VPN too. But on NuttX, a VPN can become a normal network device very naturally.

6 — How It Started: February to October (2:46 → 3:38)

On slide: A month axis Feb → Oct. Above it: mid Feb, read the GSoC idea list / early Mar, plan written up as an issue / Mar 20, heard about this conference from Alan, sent a talk / early Apr, GSoC proposal / early May, talk accepted, GSoC not selected / Oct, Glasgow. Below it: Jun–Jul a little at a time, August holiday most of the work, September FLAT app → kernel driver.

Say:

- Here is how this project went, month by month.
- In the middle of February, I read the GSoC idea list and started to look into it. In early March, I wrote my plan as an issue and proposed it. On March 20, Alan told me about this conference, so I sent a talk proposal. In early April, I sent my GSoC proposal.
- In early May, I got two answers. The talk was accepted, and the GSoC proposal was not. I decided to build it anyway.
- From June, I worked on it a little at a time. I did most of the work during the summer holiday in August. In September, I moved it from a FLAT app into a kernel driver. That is the second half of this talk.

7 — Background: What Is NuttX (3:38 → 4:26)

On slide: Software stack (applications / POSIX, ANSI, BSD socket API / networking, VFS, file systems, device drivers, NSH / NuttX kernel / MCU, MPU) and tags: RTOS, POSIX-oriented, BSD sockets, VFS, NSH, native networking, device drivers, ASF top-level project.

Say:

- First, what is NuttX?
- NuttX is a real-time operating system that puts standards first. It follows the POSIX and ANSI APIs, it uses BSD sockets for networking, and it has a real file system, a shell called NSH, built-in networking and device drivers. All of this is made small enough to run on microcontrollers, from 8-bit to 64-bit. It

is also a top-level project of the Apache Software Foundation, and its community develops it in the open.

- The one point I want you to remember: NuttX is a small RTOS, but you program it like Unix. That is rare. And that is why almost everything later in this talk — sockets, network devices, a POSIX-style port — is possible at all.

8 — Where NuttX Fits (4:26 → 5:32)

On slide: FreeRTOS / Zephyr / NuttX table (design centre, typical use), GitHub stars with the note that stars are visibility, not deployment; number cards: 15+ architectures, 300+ boards, 1500+ configs, Sony, PX4, Xiaomi OpenVela 1000+ SKUs, Japan's 2024 lunar mission.

Say:

- Thirty seconds on where NuttX fits, and whether it is a toy.
- These are three good RTOSes, and each has a different focus. FreeRTOS is a small real-time kernel, on purpose, for MCU firmware and control. Zephyr is an all-in-one platform for connected devices, wearables and industrial IoT. NuttX is the Unix-like one: networking, file systems and drivers, used in drones and sensing. This is not a ranking. By GitHub stars, NuttX is the smallest of the three, and I want to say that clearly. But stars show how many developers notice a project, not how many devices use it.
- And it is really used. Upstream, NuttX supports more than fifteen CPU architectures, more than three hundred boards, and more than fifteen hundred configurations. In real products it runs in Sony devices, PX4 flight controllers, and more than a thousand Xiaomi OpenVela products, and it flew on Japan's 2024 Moon mission. It is small next to the big ones, but it is not a toy.

Check before the talk: the star counts move. Re-check them on the day, or leave the date on the slide as it is.

9 — History: NuttX at Sony (5:32 → 6:34)

On slide: Timeline: 2015 Sony audio products / 2018–19 SPRESENSE, CXD56xx (OSS + ELC Europe 2019) / 2020 upstream collaboration (NuttX Online Workshop, "A Journey from Fork to Mainline") / 2020s sensing and edge devices.

Say:

- This is Sony's history with NuttX. It is not my story yet, and it is older than edge AI.
- Sony has sold audio products running NuttX since around 2015. A few years later, SPRESENSE and its CXD56xx chip added audio, GNSS, camera and multi-core support to NuttX. In 2019, Masayuki Ishikawa from Sony presented this work at Open Source Summit and Embedded Linux Conference Europe. After that, Sony engineers became major upstream contributors. One of them, Alin Jerpelea, later chaired the NuttX project management committee. His 2020 talk was called "A Journey from Fork to Mainline."
- I am not saying that one engineer made Sony choose NuttX. I have no source for that, so I will not say it. What public talks and documents do show is simpler: Sony used NuttX for many years before

edge AI, and over time it spread into camera, GNSS and sensing work. My own work builds on that.

10 — Experience: SPRESENSE and AITRIOS (6:34 → 7:31)

On slide: Photos of SPRESENSE and the AITRIOS edge AI camera; the flow "application development → system integration → deployment & operation".

Say:

- That was Sony's history. This is mine. Same hardware family, but I see it from a different place.
- On SPRESENSE, I use the NuttX SDK for sensing, camera, GNSS and LTE — prototypes and embedded development, on the application side. On AITRIOS edge AI cameras, which are built with ESP32 and NuttX, my job is to set up cameras and keep them running on site, not on a desk next to a debugger.
- In both cases, I come to NuttX from the same direction: first application development, then putting the system together, then setting it up and running it. I never had to write a NuttX driver for any of that — until this project. And the question that brought me here was not "How do I make NuttX run?" It was "How do I reach a NuttX device after it is installed?"

11 — Problem: Reaching a Device From Far Away (7:31 → 8:23)

On slide: Engineer → Internet → router → Wi-Fi → NuttX device, showing that WPA2/WPA3 covers only the Wi-Fi hop; cards: open it to the Internet / make your own protocol / use a vendor's cloud.

Say:

- Now, the problem itself.
- Wi-Fi security is real, but it protects only one part of the path: the radio link between the device and the access point. Everything before that — the engineer, the Internet, the router — is outside what WPA2 or WPA3 can protect. What I needed was a secure path for the whole way, from the engineer to the device, not just the last few meters.
- Without that, every choice has a problem. You can give the device a public IP and hope. You can make your own secure protocol and keep it working forever. Or you can hand the connection to a vendor's cloud. The same problem appears in edge AI, industrial IoT, and remote systems in general. It is not limited to this project.

12 — Options: Other Ways to Reach a Device (8:23 → 9:40)

On slide: A table of seven ways (open a port / HTTPS to a cloud API / MQTT broker / reverse tunnel or relay / VPN on the router / carrier closed network / VPN on the device) against: who connects, any service or not, server in the middle, where security ends, what NuttX already has.

Say:

- Before I chose WireGuard, I compared the usual ways to reach a device.
- Opening a port is easy, but then the whole Internet can reach the device. HTTPS and MQTT are what most IoT products use. The device connects out to a server, so NAT is not a problem, and NuttX

already has clients for both. But you only get the messages you designed — you cannot log in or open a web page on the device. And the security ends at the broker or the cloud.

- A reverse tunnel gives you one port through a relay server. A VPN on the router is great when the router is yours — but the last part of the path, on the LAN, is not protected, and at a customer's site the router is usually not yours. A carrier closed network works only over cellular, and you need a contract.
- A VPN on the device is the only row that has all three: any IP service, no server in the middle, and security all the way to the device. MQTT and HTTPS are still the right tools for data. This is for maintenance access.

If asked about Tailscale or ZeroTier: they are also a VPN on the device, but they usually need a bigger OS than a microcontroller has today (footnote on the slide).

13 — Solution: Why WireGuard (9:40 → 10:33)

On slide: Engineer / Linux / Windows ↔ WireGuard tunnel ↔ NuttX device; modern IP-layer VPN / UDP transport, public-key peers / already everywhere.

Say:

- So, why WireGuard?
- WireGuard is a modern VPN that works at the IP layer. It is small and simple by design, with much less code than IPsec or OpenVPN, and that matters a lot on a microcontroller. It runs over plain UDP, and each peer is known by its public key, not by certificates or a long setup. If a packet does not carry the right key, WireGuard does not answer at all — to everyone else, the device is just a brick wall. And most importantly, it is a standard. It has been in the Linux kernel for years, there is an official Windows client, and the other end can be any WireGuard peer.
- So I chose a standard protocol, instead of making my own or depending on a vendor's cloud.

14 — Architecture: WireGuard on the NuttX Network Stack (10:33 → 11:21)

On slide: Application → TCP/UDP/IP → wg0 (NET_LL_TUN netdev) → WireGuard encrypt/decrypt → UDP socket → wlan0 / usrsoc → network, with TX and RX arrows.

Say:

- This is one step more technical than the overview. This is where the implementation starts.
- Going out, a packet comes down from the application through TCP or UDP and IP, and reaches wg0 like any other network device. WireGuard encrypts it, and it leaves as a UDP packet — through wlan0 on a board, or through usrsoc when the Wi-Fi chip itself works through sockets. Coming in is the same path the other way: a UDP packet arrives, WireGuard decrypts it, and it is passed up to wg0 as if it came in on a normal interface.
- From here on, the work is building this middle part — wg0 and the encryption — without changing the layers above or below it.

15 — Agenda: ② Reusing an Existing WireGuard (11:21 → 11:31)

Say (one line): That is the background. Next: how I reused an existing WireGuard.

Pause: point at the highlighted line, one beat.

16 — Porting Strategy: Reusing wireguard-lwip Without lwIP (11:31 → 12:25)

On slide: Three cards: protocol & crypto (KEEP) / platform hooks (ADAPT, the four functions in wireguard-platform.h) / lwIP network glue (REPLACE).

Say:

- Why I started from smartalock/wireguard-lwip, and did not write the crypto myself.
- wireguard-lwip is cleanly split into three layers. The protocol and crypto — the handshake, the data transport, BLAKE2s, ChaCha20-Poly1305 and Curve25519 — are portable C that does not depend on any OS, so I used them without changes. Everything that depends on the OS was already kept behind one header, wireguard-platform.h, as four functions. I only had to replace the lwIP network glue — the part that registers an lwIP netif and handles lwIP buffers — because NuttX does not have lwIP at all.
- I did not have to rewrite the protocol or the crypto, and that was not luck. That is what a clean line between the protocol and the OS gives you. Because of it, this port took weeks, not months.

17 — Design: Mapping lwIP onto NuttX netdev (12:25 → 12:47)

On slide: netif → net_driver_s, pbuf → iob, netif_add() → netdev_register(), callbacks → devif_poll().

Say:

- This was the most important porting decision.
- Instead of bringing lwIP into NuttX, I matched each lwIP idea to NuttX's own. netif becomes net_driver_s, pbuf becomes iob, netif_add becomes netdev_register, and the send callback becomes NuttX's devif_poll. wg0 is registered as a TUN-type network device, and the “wire” under it is a UDP socket.

18 — Result: A Working Tunnel Is Only the Beginning (12:47 → 13:21)

On slide: Protocol & crypto: 3,079 lines unchanged / OS hooks: four functions / network glue: the real work. wg genkey / set / setconf at run time.

Say:

- Here is the result of the first version.
- The protocol and crypto came over exactly as they were — more than three thousand lines, unchanged. The four OS functions are small. The real work was the network glue. And it worked: you

can make a key, set peers and load a config, all at run time, with no rebuild. So far, a success story — but the princess was in another castle. The rest of this talk is about what that success was hiding.

19 — Agenda: ③ Bugs a Working Demo Hid (13:21 → 13:31)

Say (one line): That was the porting. Now, the part this whole talk is built around.

Pause: point at the highlighted line, one beat.

20 — The Testing Pattern: Simple Tests Pass, Deep Tests Fail (13:31 → 13:59)

On slide: SIMPLE TEST (handshake / ping, PASS) → ASSUMPTION ("that part must be fine") → DEEPER TEST (TCP / real hardware / kernel / security, FAIL).

Say:

- This is the main point of the talk.
- Every hard bug in this project looked the same. A simple test passes, but a deeper one fails. Even worse, the passing test makes you believe that part is fine, so you look everywhere else for hours. The next slides are examples of this pattern — five in the apps version, and two more after I moved it into the kernel.

21 — Bug 1/7: SO_RCVTIMEO (13:59 → 14:23)

Say:

- The first one: a success code that did nothing.
- Setting a receive timeout returned zero, which means success. But `recvfrom` still waited forever. The option was accepted and then silently ignored. And that made a perfect deadlock: the timer never fired, so no handshake was sent, so the peer never answered, so `recvfrom` never returned to let the timer fire.

22 — Bug 2/7: The Detached pthread (14:23 → 15:07)

Say:

- The second: something I believed about a thread's life that was not true.
- The receive thread was created as detached. When the command that started it finished, the thread stopped too — it did not even appear in `ps`. When I added a short sleep after creating it, its `printf` ran, so it really was created. My mistake was thinking that "detached" means "it lives longer than the command." Detach only decides when the thread's memory is freed, not whether it keeps running after the command exits. So when the command ended, the thread ended with it. That is why, in the kernel version, the driver itself owns the thread.

23 — Bug 3/7: ping Works, TCP Dies (15:07 → 15:59) (*my favorite*)

Say:

- This is my favorite bug in the project.
- On the same tunnel, ping was perfect, and even a TCP handshake finished — and then the connection died. When I added logging to the send path, it showed EBADF: bad file descriptor. All the sends that worked happened inside the receive task itself — handshake replies, keepalives, ping replies, the SYN-ACK. Sends from any other task failed.
- The reason: in NuttX, a file descriptor belongs to one task group. It is not shared by the whole system. The fix was to stop using file descriptors and use NuttX's internal socket API. It works on a struct socket, which is just memory, so any task can use it. That is exactly why the kernel version keeps the socket as a struct socket.

24 — Bug 4/7: "No Real Impact" Was 10× (15:59 → 16:39)

Say:

- The fourth: something I never measured.
- The fix without file descriptors also broke poll, because poll works with file descriptors. So, as a quick fix, I checked the socket every fifty milliseconds. I thought the only cost was waking up the CPU when it was idle. I was wrong. Every packet that arrived waited up to fifty milliseconds, and that alone limited the speed — about ten times slower than it should have been. I thought it had no real cost, and it cost ten times. The real fix was a proper wait: a callback on psock_poll and a semaphore.

25 — Bug 5/7: sim Passes, Real Hardware Is Different (16:39 → 17:33)

Say:

- The fifth, and it will come back later: stack sizes nobody had measured.
- Two things only real hardware showed me. First, saving the config into a file through the shell froze the board — but only on the flash file system of real hardware, never in sim. I fixed it with a wg saveconf command that writes the file itself.
- Second, the receive task's stack. That task calls into the IP stack on its own stack, so the whole TCP call chain sits on it, and for packets it answers right away, so does encrypting the reply. The default size, 3072 bytes, was a number nobody had measured. When I measured it on the board with stack coloration and ps, it was ninety-four percent full — only 176 bytes left. Please remember this one.

26 — Agenda: ④ Into the Kernel (17:33 → 17:43)

Say (one line): Those were the five bugs. Now the harder half: making this a real kernel driver.

Pause: point at the highlighted line, one beat.

27 — Part 4: Into the Kernel (17:43 → 18:18)

Say:

- This is the work I did after the first version, and it is the part that makes it ready to merge.
- Version one lives in apps and uses socket internals. That is fine for a FLAT build. But it can never be a proper in-tree driver, and it does not work in PROTECTED and KERNEL builds, where user space and the kernel are kept apart. So I moved the device into the kernel, under drivers/net/wireguard, and I control it from a small user-space wg command through ioctl.

28 — The Kernel Design (18:18 → 19:03)

Say:

- This is the shape of the real driver.
- wg0 is now a lower-half network device in the kernel. It owns its UDP socket and its receive thread, so the socket and thread problems from part three cannot happen, by design. The settings go through a simple ioctl ABI with no pointers. No pointers, because the kernel copies the caller's struct as it is. The private key is write-only: no ioctl ever gives it back — but it is still in the config file on the user side. The crypto in the kernel is NuttX's own. The only copied code is a small MIT-licensed X25519 in the user-space command, for making keys offline.

29 — Two More Bugs, Same Shape (19:03 → 20:08) (*locked with 26*)

Say:

- And the same pattern came back in the kernel — twice.
- First, a bug in NuttX's own ChaCha20-Poly1305: the counter was written into the wrong bytes of the nonce. When the counter is zero, both layouts are the same, so the handshake works, and then every data packet after the first one fails to decrypt. Simple test passes, deep test fails. Nothing in NuttX had called this code before, so nobody had hit it. It is a two-line fix, and it goes upstream as its own crypto PR. So porting to NuttX found a hidden bug in NuttX itself.
- Second, a crash that happened only in a kernel build. Setting the key copied the peers into a six-kilobyte array on the stack, but the kernel stack is only three kilobytes. Overflow, a broken heap, panic. In sim, with big task stacks, it never showed up. It was the same stack lesson as before, one level deeper. I moved the array to the heap.

30 — Agenda: ⑤ Checking That It Really Works (20:08 → 20:18)

Say (one line): The driver is done. Now: how do I know it really works?

Pause: point at the highlighted line, one beat.

31 — Testing: Seven Places It Must Work (20:18 → 24:32) (*core — do not cut*)

On slide: Table (where / version / what I checked): sim / rv-virt knetnsh64 (BUILD_KERNEL) / rv-virt pnsh64 (BUILD_PROTECTED) / rv-virt knetnsh64_smp (4 CPUs) / SPRESENSE measured / ESP32-S3 + SPRESENSE over real Wi-Fi / apps v0.1.1 on both boards.

Say:

- Let me be exact about which version was checked where.
- The kernel driver builds with no errors in the sim config — the config that CI builds — and it runs as a full tunnel in a real kernel build: rv-virt knetnsh64 in QEMU, over virtio-net to WireGuard in the Linux kernel. The wg command is loaded as a separate program, and it reaches the driver through system calls.
- Next, the other build type. The ioctl ABI is a fixed struct with no pointers inside, so the kernel never has to read user memory. If that idea is wrong, you do not get a compile error. You get an MPU fault on the first call. PROTECTED had only been built before, on a board with no emulator. So I ran it on rv-virt's protected config: it boots with the memory split, the wg command lives on the user side, every ioctl crosses the boundary, and the tunnel reaches Linux WireGuard. But when I ran it again and again, two runs out of eight hit a fault in user space. I have not found the cause yet, so I list it as an open problem, not as a pass. The memory map is the default one — the kernel uses eighty percent of its part, and the user command uses eight percent of the other.
- Then, locking. The driver protects its protocol state with the device lock. It does not depend on there being only one CPU. Until late September, I could only argue that from the code; I had never really tested it. So: four CPUs, two thousand packets coming in, while another CPU changes the settings and takes the interface down and up. No fault, no assertion, and the tunnel came back. One scripted test on an emulator is not a full search for race conditions, so a race that needs very exact timing could still be hiding. But the claim is no longer only on paper.
- The last row is the one I would show a reviewer first. The long test in sim showed that memory came back exactly to where it started. That is good, but the simulator uses the host's memory allocator and 64-bit frames, so none of those numbers apply to a board. So I measured the board itself: a hundred and twenty seconds of traffic, then forty down-and-up cycles. Each one closes and reopens a socket on the Wi-Fi module through usrsock and does a new handshake over real Wi-Fi. Heap, number of allocations, buffer pool, and the peak of every stack: exactly the same before and after. And the number I really wanted — the most stack the receive thread ever used — is one thousand four hundred and seventy-two bytes, out of a six-kilobyte default. That default used to be a guess. Now it is a measurement.
- And now it runs on real hardware — both boards. I flashed the kernel driver to an ESP32-S3, with a tunnel to the official Windows WireGuard client over real Wi-Fi. I also flashed it to a SPRESENSE. Its GS2200M Wi-Fi works through usrsock, so the driver really sends through usrsock on real hardware. Then I forced the exact failure this design is for, on that board: a packet stuck waiting on the real usrsock side while control commands arrive. A query got an answer, a peer update was accepted, the stop command reported a timeout instead of hanging, the stuck packet was left alone, and a second down cleaned it up, so the interface came back. This still does not cover running out of IOBs, losing

Wi-Fi in the middle of a send, or long heavy traffic. The older apps version on both boards — telnet, HTTP, a seven-megabyte transfer, rekeying, and recovery after power loss — was also checked.

32 — Testing: A Working Tunnel Does Not Prove Enough (24:32 → 27:09) (core — do not cut)

On slide: Three rows (same private key every boot / session keys that outlive the tunnel / a leak that shows after hours), each with "would ping notice?" → no.

Say:

- The last slide was about where it runs. This one is about what a working tunnel cannot tell you.
- Take making keys. NuttX can provide /dev/urandom with an xorshift128 backend, and its seed can come from build-time settings. A board left like that can give out the same private key every time it starts. The handshake works. Ping works. The demo works. But the tunnel is worthless, and nothing I have shown you so far would notice. So I reset the board three times and compared one new key from each start: three different keys, with the entropy pool backend chosen, not xorshift128. That is a useful check that the problem has not come back, but it does not prove the randomness is good.
- Second, what is left after you take the interface down. Here I stopped reading the code and read the memory instead. I found the device in the running system by searching for the key I had set, and then checked each secret field by name, at offsets taken from the driver's own headers. The transport keys are filled while the tunnel is up, and every byte is zero after down. The handshake data is already gone while it is up, because it is erased as soon as the keys are made. The static key is still there — on purpose, because up has to work again — and there is exactly one copy of it in all writable memory, before and after. This difference is a decision, so it is now written in the driver's documentation, and a reviewer does not have to guess.
- Third, the long test. A hundred and fifty-one forced rekeys, past the point where sessions expire. Ten endpoint changes. A hundred down-and-up cycles. And memory checked the whole time. The buffer pool comes back complete, and the number of live allocations goes back to exactly the number at start-up — not close, the same. The receive thread's stack peak is two and a half kilobytes, and it does not grow after the first sample. That last number is useful by itself: the default stack for that thread is six kilobytes, so the safety margin is real, not just hoped for. What I am not saying is that it ran for a week, or that the simulator's allocator and 64-bit frames act like the board.

33 — A Design Choice: The Timestamp Problem Has No Easy Answer (27:09 → 30:57) (core — do not cut)

On slide: What replay protection obliges, the four options (real-time clock + in-boot high-water mark / persist every time / durable range reservation / wait to catch up), and the measured pair: clock unset → no handshake in 75 s, clock set → 4.1 s.

Say:

- This is the one place where I had to decide, not test. And then I measured my decision.

- WireGuard's replay protection gives the side that starts the handshake a rule: every handshake must have a newer timestamp than anything that peer has already accepted — and this must stay true after a reboot. The device is usually the side that starts, and it is also the side least likely to have a clock with a battery.
- In my simulator test with the old uptime-based timestamp, recovery took seventy-five seconds, but that number is misleading. The device has to get past the last handshake of its previous session, so it is cut off for as long as it ran last time. It was seventy-five seconds only because that session was short. A device that ran for a month cannot reconnect for a month. So doing nothing is not an option.
- Saving the timestamp at every handshake means writing to flash every few minutes, forever, on the path that most needs to be fast. Reserving a range in flash and using it up from RAM is the right long-term answer, and I have written up that design. But it needs storage that promises the order and safety of its writes, and I could not get that. On my machine, hostfs sync calls an empty function and returns OK, so a successful write in the simulator proves nothing about power loss. Shipping that would look solved without being solved. Here, the easy choice and the right choice were not the same.
- So I chose the real-time clock plus the highest value used in this boot. That is exactly what Linux and wireguard-go do, and it never falls back to uptime. What I added is the part a desktop never needed, because a desktop always has a clock. Without an RTC, NuttX starts the clock from a set year. So if the clock is still before the next year, it was never set — and the driver prints a warning once, and says what will happen. It still sends the timestamp, because a peer that meets this key for the first time will accept it, and refusing would break first pairing on exactly those boards. This one line of output is for a failure that does not crash: a tunnel that silently never comes back, on a device that is not on your desk.
- And I did measure it on the board. The trap is that the obvious test passes. Reset the board, ping it, and the tunnel is back in sixteen seconds — because your ping made the peer start the handshake, and the side that answers does not need its own timestamp. The device could be completely unable to connect out, and that test still passes.
- So I sent the board nothing at all, and watched how it saw its own session. Two runs, same board, same peer, same image, only one change. With the clock as it boots, the driver warns, and there is no handshake after seventy-five seconds. With the clock set from my laptop, the warning stops, and the handshake finishes in four point one seconds. That shows both sides at once: the unsolved case is real on hardware, not just in a simulator, and the design works wherever there is a clock. Issue fourteen stays open because I decided it should, not because I ran out of evidence.

34 — Agenda: ⑥ Demo (30:57 → 31:07)

Say (one line): That was all engineering. Now, what it actually looks like.

Pause: point at the highlighted line, one beat.

35 — Demo Setup: One Tunnel on a Phone's Network (31:07 → 32:05)

On slide: Diagram: venue Wi-Fi (blocks devices from talking to each other) → phone hotspot (its own small local network) → laptop (WireGuard peer in Docker, wg0 10.10.0.1; voice clips on :8000) and StackChan (ESP32-S3 + NuttX, wg0 10.10.0.2 over wlan0 on the hotspot), joined by one WireGuard tunnel; SPRESENSE as a dashed second peer (wg0 10.11.0.2).

Say:

- Here is the demo setup.
- My laptop and the StackChan are both on my phone's hotspot. The venue Wi-Fi works, but it blocks devices from talking to each other, so the phone makes a small local network for us. The phone can use the venue Wi-Fi or mobile data for its own Internet, but the demo does not need it.
- Between the laptop and the StackChan there is one WireGuard tunnel. The laptop side is a normal Linux WireGuard in a container, at 10.10.0.1. The StackChan is 10.10.0.2. Everything in the demo — telnet, the web page, and even the voice — goes through that tunnel. I made the voice clips on the laptop in advance, so there is no cloud API anywhere.
- Please watch two things: the addresses after I log in, and the packets. Outside the tunnel, you only see encrypted UDP.

36 — The Demo (32:05 → 33:35)

On slide: Seven steps: wg show / telnet 10.10.0.2 / ifconfig / stackchan face happy / stackchan say hello.wav / tcpdump outside vs. inside / wg show again. Right panel: outside = only encrypted UDP 51820, inside = the telnet text.

Say:

- The point of this demo is that there is nothing special to see. The StackChan is just a normal computer at the far end of an encrypted tunnel.
- First, wg show on the laptop: the handshake is done, and you can see the StackChan's address on the phone's network. Then I telnet to 10.10.0.2 — that is the address inside the tunnel. When I type ifconfig, you see both: wlan0 on the phone's network, and wg0 inside the tunnel. So I am really logged in through the tunnel.
- Now I can control it. Let's make it happy, and let it say hello. The voice comes from the laptop through the same tunnel.
- Last, the packets. On the left, the phone's network: only encrypted UDP. On the right, inside the tunnel: the same telnet session, in plain text. And wg show again: the byte counts went up. Mischief managed.

Live: scripts/stackchan/demo.sh: login in the left terminal, outside and inside in two terminals on the right, show before and after. Put the terminals on the left and the StackChan (camera or on the desk) on the right.

Before the talk: phone hotspot on (2.4 GHz / maximize compatibility), laptop on it, Docker Desktop running, the StackChan joined with nsh_wifi.py (the passphrase is not saved on the board), wg_setup.py

run once. Check `stackchan say` once. If the hotspot was turned off and on, its addresses may change: run `nsh_wifi.py` and `wg_setup.py` again. Rarely (about 1 in 20) the board stops during `say`; power it off and on, then repeat the same steps.

Never on screen: `wg showconf` (prints the private key), the key files in `%USERPROFILE%\stackchan-wg`, the hotspot passphrase.

Fallback: the recording. Switch at once if the live demo stalls.

37 — Agenda: ⑦ What Comes Next (33:35 → 33:45)

Say (one line): Last part: going from a demo to something you can run, and what to take home.

Pause: point at the highlighted line, one beat.

38 — Running It for Real (33:45 → 34:19) (*cut candidate 3*)

Say:

- Going from “it works in a demo” to “you can really run it.”
- Keys are no longer built into the firmware. You make them and set them at run time, and they never end up in the firmware image. And for long runs, I check four separate signals every minute — is the router side alive, is the board on the network, is WireGuard carrying traffic, is the application data moving — so when it goes quiet, I can say which part failed instead of guessing.

39 — Other CPUs (34:19 → 34:46) (*cut candidate 2*)

Say:

- Did the decision from the start work?
- Keeping the OS behind a few functions was a bet. Here is the check, with the apps FLAT version: four CPU types — x86, ARM Cortex-A7, Xtensa on the ESP32, and Cortex-M4F on SPRESENSE — and no code change to move between them. The kernel driver also runs on sim and on the rv-virt kernel build. The bet paid off.

40 — Giving Back (34:46 → 35:29) (*cut first*)

Say:

- Let me end this part with the values of this community.
- I will send the work in steps. First, a small crypto PR for the nonce fix, because it is useful on its own and helps everyone. Then the driver PR for the kernel device, and then the command PR for the user-space tool. The work is already split into small PRs, and every commit passes the patch checks and builds on its own. The design discussion and the upstream PRs are the next step; they have not started yet. Community over code means being honest about where the work really is, and then letting review change it.

If asked "is it merged?": not yet. These are submission candidates, not a claim that it is mergeable. See the Q&A file.

If time is short: skip the whole slide. The talk's argument (the pattern, the verification) does not rest on it; just keep the answer above ready for questions.

41 — Takeaways (35:29 → 36:31)

Say:

- Three things to take home.
- First, simple tests pass and deep tests fail — so design your tests to go deep: long data runs, real hardware, a real kernel build, not just a handshake. Second, a clean line between your code and the OS is worth the effort — four CPU types and three NuttX versions, with no changes. Third, porting to a project is a chance to make it better. One bug was in NuttX's own crypto, and that fix goes upstream by itself. I found the second only because I wanted my board image to build from a clean checkout: a normal release tag does not boot to the shell on this board, because before the RTC starts, the clock reads zero, and the watchdog that should finish starting the RTC never fires. And the third was in my own driver, and only a real kernel build showed it; sim never would have.

42 — Thank You (36:31 → 36:40)

Say:

- Thank you. I am happy to take your questions — about the kernel driver, the bugs, the ioctl ABI, or the hardware.

After: keep [coc-glasgow-qa.md](#) at hand, not a slide. The usual questions (rebase status, #14, IOB exhaustion, checkpatch) are answered there and on no slide.